

Assignment 06

Name: Kesav Adithya Venkidusamy

Course: DSC650 - Big Data

Instructor: Amirfarrokh Iranitalab

Assignment 6.1

Using section 5.1 in Deep Learning with Python as a guide (listing 5.3 in particular), create a ConvNet model that classifies images in the MNIST digit dataset. Save the model, predictions, metrics, and validation plots in the dsc650/assignments/assignment06/results directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

```
In [1]: # Load all the required libraries
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt
from keras.utils import to_categorical
from keras import models
from matplotlib import pyplot
from keras.preprocessing.image import ImageDataGenerator
import numpy as np
from keras.optimizers import SGD, Adam
```

Load the MNIST dataset

```
In [15]: # Load the MNIST dataset
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28, 28, 1))
train_images = train_images.astype('float32') / 255
test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

Model Building

```
In [16]: ## Instantiating a convnet
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
```

```
In [17]: ## Adding classifier on top of convnet
model.add(layers.Flatten())
model.add(layers.Dense(10, activation='softmax'))
```

```
In [18]: ## Showing model summary
model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 26, 26, 32)	320

max_pooling2d_4 (MaxPooling2D)	(None, 13, 13, 32)	0

conv2d_9 (Conv2D)	(None, 11, 11, 64)	18496

max_pooling2d_5 (MaxPooling2D)	(None, 5, 5, 64)	0

conv2d_10 (Conv2D)	(None, 3, 3, 128)	73856

flatten_1 (Flatten)	(None, 1152)	0

dense_2 (Dense)	(None, 10)	11530
=====		
Total params: 104,202		
Trainable params: 104,202		
Non-trainable params: 0		

```
In [19]: # Compile the model
model.compile(optimizer="rmsprop",
              loss="categorical_crossentropy",
              metrics=["accuracy"])
```

Model Validation

```
In [20]: ## Set aside a validation set (10000 samples)
# Data
validation_images = train_images[:10000]
partial_train_images = train_images[10000:]
# Labels
validation_labels = train_labels[:10000]
partial_train_labels = train_labels[10000:]
```

Model Training



```
In [21]: # Train the model
history = model.fit(partial_train_images,
                    partial_train_labels,
                    epochs=5,
                    batch_size=64,
                    validation_data=(validation_images, validation_labels))
```

```
Epoch 1/5
782/782 [=====] - 11s 15ms/step - loss: 0.1793 - accuracy: 0.9442 - val_loss: 0.1548 - val_accuracy: 0.9484
Epoch 2/5
782/782 [=====] - 11s 14ms/step - loss: 0.0491 - accuracy: 0.9849 - val_loss: 0.0538 - val_accuracy: 0.9850
Epoch 3/5
782/782 [=====] - 11s 14ms/step - loss: 0.0339 - accuracy: 0.9893 - val_loss: 0.0475 - val_accuracy: 0.9868
Epoch 4/5
782/782 [=====] - 11s 14ms/step - loss: 0.0252 - accuracy: 0.9921 - val_loss: 0.0511 - val_accuracy: 0.9868
Epoch 5/5
782/782 [=====] - 11s 14ms/step - loss: 0.0193 - accuracy: 0.9939 - val_loss: 0.0400 - val_accuracy: 0.9897
```

Plotting Model Output and Loss

```
In [22]: # Plot the training and validation accuracy and loss
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
```

```
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
```



Evaluate the Model

```
In [23]: # Evaluate the convnet
test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f'Test accuracy: {test_acc*100:.1f}%')
print(f'Test loss: {test_loss:.3f}')
```

```
313/313 [=====] - 1s 4ms/step - loss: 0.0289 - accuracy: 0.9911
Test accuracy: 99.1%
Test loss: 0.029
```

The model accuracy is 99.1% and loss is only 0.029; The accuracy is increased significantly and loss is reduced a lot by adding Conv2D and MaxPooling2D layers

Save Model

```
In [24]: model.save('results/mnist')
```

```
INFO:tensorflow:Assets written to: results/mnist/assets
```

Assignment 6.2

Assignment 6.2.a

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies images CIFAR10 small images classification dataset. Do not use dropout or data-augmentation in this part. Save the model, predictions, metrics, and validation plots in the dsc650/assignments/assignment06/results directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

Load the data & Data preparation

```
In [25]: # Load the CIFAR10 data set
(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()
assert x_train.shape == (50000, 32, 32, 3)
assert x_test.shape == (10000, 32, 32, 3)
assert y_train.shape == (50000, 1)
assert y_test.shape == (10000, 1)
```

```
In [26]: # summarize loaded dataset
print('Train: X=%s, y=%s' % (x_train.shape, y_train.shape))
```


```
print('Test: X=%s, y=%s' % (x_test.shape, y_test.shape))
```

Train: X=(50000, 32, 32, 3), y=(50000, 1)

Test: X=(10000, 32, 32, 3), y=(10000, 1)

In [27]:

```
# plot first few images
for i in range(9):
    # define subplot
    pyplot.subplot(330 + 1 + i)
    # plot raw pixel data
    pyplot.imshow(x_train[i])
# show the figure
pyplot.show()
```



In [28]:

```
## Set aside a validation set (10,000 samples)
# Data
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
# Labels
y_val = y_train[:10000]
partial_y_train = y_train[10000:]
```

Model Building without dropout or data-augmentation

In [29]:

```
## Instantiating a convnet
model = models.Sequential()
```

```

model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)))
model.add(layers.experimental.preprocessing.Rescaling(1./255))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(128, (3, 3), activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))

```

```

In [30]: ## Adding classifier on top of convnet
model.add(layers.Flatten())
model.add(layers.Dense(10, activation='softmax'))

```

```

In [31]: ## Showing model summary
model.summary()

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_11 (Conv2D)	(None, 30, 30, 32)	896
rescaling (Rescaling)	(None, 30, 30, 32)	0
max_pooling2d_6 (MaxPooling2	(None, 15, 15, 32)	0
conv2d_12 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_7 (MaxPooling2	(None, 6, 6, 64)	0
conv2d_13 (Conv2D)	(None, 4, 4, 128)	73856
max_pooling2d_8 (MaxPooling2	(None, 2, 2, 128)	0
flatten_2 (Flatten)	(None, 512)	0
dense_3 (Dense)	(None, 10)	5130
=====		
Total params: 98,378		
Trainable params: 98,378		
Non-trainable params: 0		

```

In [32]: # Compile model
model.compile(optimizer="rmsprop",

```

```
loss="sparse_categorical_crossentropy",
metrics=["accuracy"]])
```

Model Training

In [33]:

```
# Train model
history = model.fit(partial_x_train,
                    partial_y_train,
                    epochs=30,
                    batch_size=64,
                    validation_data=(x_val, y_val))
```


```
Epoch 1/30
625/625 [=====] - 13s 20ms/step - loss: 1.6908 - accuracy: 0.3901 - val_loss: 1.4259 - val_accu
acy: 0.4851
Epoch 2/30
625/625 [=====] - 12s 20ms/step - loss: 1.3411 - accuracy: 0.5253 - val_loss: 1.1883 - val_accu
acy: 0.5744
Epoch 3/30
625/625 [=====] - 12s 20ms/step - loss: 1.1757 - accuracy: 0.5865 - val_loss: 1.0877 - val_accu
acy: 0.6131
Epoch 4/30
625/625 [=====] - 12s 20ms/step - loss: 1.0576 - accuracy: 0.6312 - val_loss: 1.0978 - val_accu
acy: 0.6277
Epoch 5/30
625/625 [=====] - 13s 20ms/step - loss: 0.9626 - accuracy: 0.6679 - val_loss: 1.0099 - val_accu
acy: 0.6548
Epoch 6/30
625/625 [=====] - 12s 20ms/step - loss: 0.8841 - accuracy: 0.6920 - val_loss: 0.9303 - val_accu
acy: 0.6802
Epoch 7/30
625/625 [=====] - 12s 20ms/step - loss: 0.8182 - accuracy: 0.7183 - val_loss: 0.9346 - val_accu
acy: 0.6833
Epoch 8/30
625/625 [=====] - 12s 20ms/step - loss: 0.7581 - accuracy: 0.7387 - val_loss: 0.9623 - val_accu
acy: 0.6784
Epoch 9/30
625/625 [=====] - 12s 20ms/step - loss: 0.7045 - accuracy: 0.7583 - val_loss: 0.9184 - val_accu
acy: 0.6901
Epoch 10/30
625/625 [=====] - 12s 20ms/step - loss: 0.6561 - accuracy: 0.7736 - val_loss: 0.8863 - val_accu
acy: 0.7017
Epoch 11/30
625/625 [=====] - 12s 20ms/step - loss: 0.6121 - accuracy: 0.7880 - val_loss: 0.8783 - val_accu
acy: 0.7119
Epoch 12/30
```



625/625 [=====] - 12s 20ms/step - loss: 0.5683 - accuracy: 0.8031 - val_loss: 1.0025 - val_accuracy: 0.6857
Epoch 13/30
625/625 [=====] - 12s 20ms/step - loss: 0.5273 - accuracy: 0.8168 - val_loss: 0.9672 - val_accuracy: 0.6940
Epoch 14/30
625/625 [=====] - 12s 20ms/step - loss: 0.4920 - accuracy: 0.8298 - val_loss: 0.9277 - val_accuracy: 0.7085
Epoch 15/30
625/625 [=====] - 12s 20ms/step - loss: 0.4547 - accuracy: 0.8434 - val_loss: 0.9808 - val_accuracy: 0.7070
Epoch 16/30
625/625 [=====] - 12s 20ms/step - loss: 0.4191 - accuracy: 0.8527 - val_loss: 1.0156 - val_accuracy: 0.7030
Epoch 17/30
625/625 [=====] - 12s 20ms/step - loss: 0.3856 - accuracy: 0.8652 - val_loss: 1.0040 - val_accuracy: 0.7049
Epoch 18/30
625/625 [=====] - 12s 20ms/step - loss: 0.3560 - accuracy: 0.8767 - val_loss: 1.1285 - val_accuracy: 0.6891
Epoch 19/30
625/625 [=====] - 12s 20ms/step - loss: 0.3272 - accuracy: 0.8863 - val_loss: 1.1621 - val_accuracy: 0.6923
Epoch 20/30
625/625 [=====] - 12s 20ms/step - loss: 0.3010 - accuracy: 0.8951 - val_loss: 1.1480 - val_accuracy: 0.7021
Epoch 21/30
625/625 [=====] - 12s 20ms/step - loss: 0.2758 - accuracy: 0.9035 - val_loss: 1.2192 - val_accuracy: 0.6969
Epoch 22/30
625/625 [=====] - 12s 20ms/step - loss: 0.2535 - accuracy: 0.9124 - val_loss: 1.2460 - val_accuracy: 0.7001
Epoch 23/30
625/625 [=====] - 12s 20ms/step - loss: 0.2339 - accuracy: 0.9182 - val_loss: 1.3402 - val_accuracy: 0.7000
Epoch 24/30
625/625 [=====] - 12s 20ms/step - loss: 0.2127 - accuracy: 0.9255 - val_loss: 1.3950 - val_accuracy: 0.6922
Epoch 25/30
625/625 [=====] - 12s 20ms/step - loss: 0.1943 - accuracy: 0.9317 - val_loss: 1.4656 - val_accuracy: 0.6919
Epoch 26/30
625/625 [=====] - 12s 20ms/step - loss: 0.1786 - accuracy: 0.9378 - val_loss: 1.4800 - val_accuracy: 0.6988
Epoch 27/30
625/625 [=====] - 12s 20ms/step - loss: 0.1652 - accuracy: 0.9421 - val_loss: 1.6897 - val_accuracy: 0.6787
Epoch 28/30
625/625 [=====] - 12s 20ms/step - loss: 0.1538 - accuracy: 0.9460 - val_loss: 1.6737 - val_accuracy: 0.6787

```
acy: 0.6943
Epoch 29/30
625/625 [=====] - 12s 20ms/step - loss: 0.1429 - accuracy: 0.9496 - val_loss: 1.7823 - val_accu
acy: 0.6849
Epoch 30/30
625/625 [=====] - 12s 20ms/step - loss: 0.1312 - accuracy: 0.9521 - val_loss: 1.8399 - val_accu
acy: 0.6826
```

Plotting the result

```
In [34]: # Plot the training and validation accuracy and loss
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
```





Evaluate the Model

In [35]:

```
# Evaluate the convnet
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f'Test accuracy: {test_acc*100:.1f}%')
print(f'Test loss: {test_loss:.3f}')
```

```
313/313 [=====] - 2s 5ms/step - loss: 1.8304 - accuracy: 0.6820
Test accuracy: 68.2%
Test loss: 1.830
```

The accuracy score without dropout and data augmentation turned out as 68.2% and loss is 1.83

Save Model

In [36]:

```
model.save('results/without_dropout_augmentation')
```

```
INFO:tensorflow:Assets written to: results/without_dropout_augmentation/assets
```

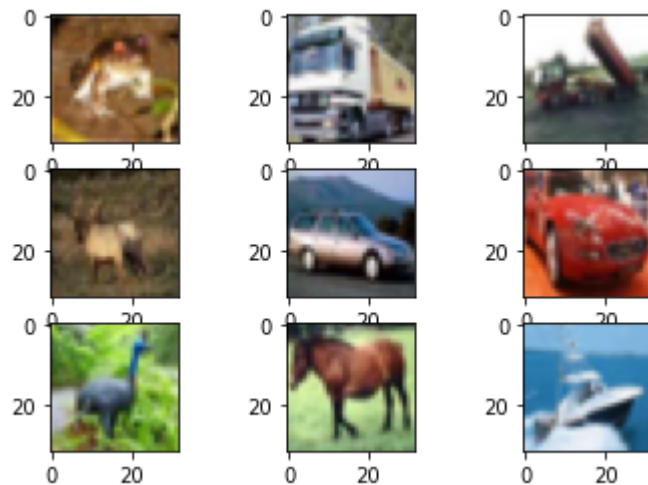
Assignment 6.2.b

Using section 5.2 in Deep Learning with Python as a guide, create a ConvNet model that classifies images CIFAR10 small images classification dataset. This time includes dropout and data-augmentation. Save the model, predictions, metrics, and validation plots in the dsc650/assignments/assignment06/results directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

Model Building with dropout or data-augmentation

In [2]:

```
#Load data
(x_train, y_train), (x_test, y_test) = keras.datasets.cifar10.load_data()
img_rows, img_cols, channels = 32, 32, 3
for i in range(0, 9):
    plt.subplot(330 + 1 + i)
    plt.imshow(x_train[i])
plt.show()
```



In [3]:

```
# set up image augmentation
datagen = ImageDataGenerator(
    rotation_range=15,
    horizontal_flip=True,
    width_shift_range=0.1,
    height_shift_range=0.1
    #zoom_range=0.3
)
datagen.fit(x_train)
```

In [4]:

```
# see example augmentation images
for X_batch, y_batch in datagen.flow(x_train, y_train, batch_size=9):
    for i in range(0, 9):
        plt.subplot(330 + 1 + i)
        plt.imshow(X_batch[i].astype(np.uint8))
    plt.show()
    break
```



In [5]:

```
#reshape into images
x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, channels)
x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, channels)
input_shape = (img_rows, img_cols, 1)
print('x_train shape:', x_train.shape)
print(x_train.shape, 'train samples')
print(x_test.shape, 'test samples')
print(y_train.shape, 'target train samples')
print(y_test.shape, 'target test samples')
```

```
x_train shape: (50000, 32, 32, 3)
(50000, 32, 32, 3) train samples
(10000, 32, 32, 3) test samples
(50000, 1) target train samples
(10000, 1) target test samples
```

In [6]:

```
#convert integers to float; normalise and center the mean
x_train=x_train.astype("float32")
x_test=x_test.astype("float32")
mean=np.mean(x_train)
std=np.std(x_train)
x_test=(x_test-mean)/std
x_train=(x_train-mean)/std
```

In [7]:

```
# Labels
```

```
num_classes=10
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

Model Building

In [9]:

```
# Building model with droupout

adm2=Adam(lr=0.001,decay=0, beta_1=0.9, beta_2=0.999, epsilon=1e-08)
opt2=adm2

model = models.Sequential()

model.add(layers.Conv2D(32, (3, 3), activation='relu', kernel_regularizer=None, input_shape=(img_rows, img_cols, channels)
model.add(layers.BatchNormalization(axis=-1))
model.add(layers.Conv2D(32, (3, 3), activation='relu',kernel_regularizer=None,padding='same'))
model.add(layers.BatchNormalization(axis=-1))
model.add(layers.MaxPooling2D(pool_size=(2, 2))) # reduces to 16x16x3x32
model.add(layers.Dropout(0.5))

model.add(layers.Conv2D(64, (3, 3), activation='relu',kernel_regularizer=None,padding='same'))
model.add(layers.BatchNormalization(axis=-1))
model.add(layers.Conv2D(64, (3, 3), activation='relu',kernel_regularizer=None,padding='same'))
model.add(layers.BatchNormalization(axis=-1))
model.add(layers.MaxPooling2D(pool_size=(2, 2))) # reduces to 8x8x3x(2*32)
model.add(layers.Dropout(0.5))

model.add(layers.Conv2D(128, (3, 3), activation='relu',kernel_regularizer=None,padding='same'))
model.add(layers.BatchNormalization(axis=-1))
model.add(layers.Conv2D(128, (3, 3), activation='relu',kernel_regularizer=None,padding='same'))
model.add(layers.BatchNormalization(axis=-1))
model.add(layers.MaxPooling2D(pool_size=(2, 2))) # reduces to 4x4x3x(4*32)
model.add(layers.Dropout(0.5))

model.add(layers.Flatten())
model.add(layers.Dense(512, activation='relu',kernel_regularizer=None))
model.add(layers.BatchNormalization())
model.add(layers.Dropout(0.5))
model.add(layers.Dense(10, activation='softmax'))

model.compile(loss='categorical_crossentropy', metrics=['accuracy'],optimizer=opt2)
```

In [10]:

```
## printing the model summary
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_2 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_3 (Conv2D)	(None, 32, 32, 32)	9248
batch_normalization_3 (Batch Normalization)	(None, 32, 32, 32)	128
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 32)	0
dropout_1 (Dropout)	(None, 16, 16, 32)	0
conv2d_4 (Conv2D)	(None, 16, 16, 64)	18496
batch_normalization_4 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_5 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_5 (Batch Normalization)	(None, 16, 16, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_2 (Dropout)	(None, 8, 8, 64)	0
conv2d_6 (Conv2D)	(None, 8, 8, 128)	73856
batch_normalization_6 (Batch Normalization)	(None, 8, 8, 128)	512
conv2d_7 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_7 (Batch Normalization)	(None, 8, 8, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 4, 4, 128)	0
dropout_3 (Dropout)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 512)	1049088
batch_normalization_8 (Batch Normalization)	(None, 512)	2048

dropout_4 (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 10)	5130
=====		
Total params: 1,345,066		
Trainable params: 1,343,146		
Non-trainable params: 1,920		

Train the model

```
In [11]: # train with image augmentation
history=model.fit(datagen.flow(x_train, y_train, batch_size=128),
                  steps_per_epoch = len(x_train) / 128, epochs=30, validation_data=(x_test, y_test))
```

```
Epoch 1/30
391/390 [=====] - 64s 165ms/step - loss: 1.9764 - accuracy: 0.3486 - val_loss: 2.3218 - val_accu
racy: 0.2912
Epoch 2/30
391/390 [=====] - 62s 160ms/step - loss: 1.4459 - accuracy: 0.4767 - val_loss: 1.4339 - val_accu
racy: 0.4932
Epoch 3/30
391/390 [=====] - 62s 159ms/step - loss: 1.2586 - accuracy: 0.5454 - val_loss: 1.5103 - val_accu
racy: 0.5330
Epoch 4/30
391/390 [=====] - 61s 157ms/step - loss: 1.1504 - accuracy: 0.5864 - val_loss: 1.2886 - val_accu
racy: 0.5760
Epoch 5/30
391/390 [=====] - 62s 158ms/step - loss: 1.0675 - accuracy: 0.6179 - val_loss: 1.2597 - val_accu
racy: 0.5807
Epoch 6/30
391/390 [=====] - 62s 158ms/step - loss: 1.0110 - accuracy: 0.6384 - val_loss: 1.1089 - val_accu
racy: 0.6319
Epoch 7/30
391/390 [=====] - 62s 157ms/step - loss: 0.9549 - accuracy: 0.6612 - val_loss: 0.9950 - val_accu
racy: 0.6728
Epoch 8/30
391/390 [=====] - 61s 156ms/step - loss: 0.9154 - accuracy: 0.6763 - val_loss: 0.9506 - val_accu
racy: 0.6776
Epoch 9/30
391/390 [=====] - 61s 156ms/step - loss: 0.8755 - accuracy: 0.6913 - val_loss: 0.9720 - val_accu
racy: 0.6753
Epoch 10/30
391/390 [=====] - 62s 157ms/step - loss: 0.8435 - accuracy: 0.7037 - val_loss: 0.8679 - val_accu
racy: 0.7084
Epoch 11/30
```





```
391/390 [=====] - 62s 158ms/step - loss: 0.8139 - accuracy: 0.7146 - val_loss: 0.7724 - val_accu
racy: 0.7394
Epoch 12/30
391/390 [=====] - 62s 158ms/step - loss: 0.7911 - accuracy: 0.7250 - val_loss: 0.7326 - val_accu
racy: 0.7519
Epoch 13/30
391/390 [=====] - 62s 158ms/step - loss: 0.7665 - accuracy: 0.7310 - val_loss: 0.8034 - val_accu
racy: 0.7368
Epoch 14/30
391/390 [=====] - 62s 159ms/step - loss: 0.7459 - accuracy: 0.7398 - val_loss: 0.7303 - val_accu
racy: 0.7541
Epoch 15/30
391/390 [=====] - 62s 158ms/step - loss: 0.7355 - accuracy: 0.7446 - val_loss: 0.8183 - val_accu
racy: 0.7331
Epoch 16/30
391/390 [=====] - 61s 157ms/step - loss: 0.7223 - accuracy: 0.7500 - val_loss: 0.7204 - val_accu
racy: 0.7611
Epoch 17/30
391/390 [=====] - 62s 159ms/step - loss: 0.7099 - accuracy: 0.7514 - val_loss: 0.6449 - val_accu
racy: 0.7830
Epoch 18/30
391/390 [=====] - 62s 158ms/step - loss: 0.6906 - accuracy: 0.7607 - val_loss: 0.7473 - val_accu
racy: 0.7574
Epoch 19/30
391/390 [=====] - 61s 156ms/step - loss: 0.6763 - accuracy: 0.7658 - val_loss: 0.7326 - val_accu
racy: 0.7666
Epoch 20/30
391/390 [=====] - 61s 157ms/step - loss: 0.6743 - accuracy: 0.7653 - val_loss: 0.7045 - val_accu
racy: 0.7646
Epoch 21/30
391/390 [=====] - 62s 158ms/step - loss: 0.6594 - accuracy: 0.7718 - val_loss: 0.6221 - val_accu
racy: 0.7891
Epoch 22/30
391/390 [=====] - 62s 158ms/step - loss: 0.6507 - accuracy: 0.7728 - val_loss: 0.6501 - val_accu
racy: 0.7822
Epoch 23/30
391/390 [=====] - 62s 158ms/step - loss: 0.6462 - accuracy: 0.7744 - val_loss: 0.6350 - val_accu
racy: 0.7875
Epoch 24/30
391/390 [=====] - 62s 158ms/step - loss: 0.6419 - accuracy: 0.7748 - val_loss: 0.6372 - val_accu
racy: 0.7934
Epoch 25/30
391/390 [=====] - 61s 156ms/step - loss: 0.6317 - accuracy: 0.7816 - val_loss: 0.6073 - val_accu
racy: 0.7954
Epoch 26/30
391/390 [=====] - 61s 156ms/step - loss: 0.6246 - accuracy: 0.7816 - val_loss: 0.6120 - val_accu
racy: 0.7960
Epoch 27/30
391/390 [=====] - 61s 157ms/step - loss: 0.6081 - accuracy: 0.7886 - val_loss: 0.5805 - val_accu
```

```
racy: 0.8099
Epoch 28/30
391/390 [=====] - 61s 156ms/step - loss: 0.6075 - accuracy: 0.7881 - val_loss: 0.6088 - val_accu
racy: 0.7971
Epoch 29/30
391/390 [=====] - 61s 156ms/step - loss: 0.5978 - accuracy: 0.7928 - val_loss: 0.6052 - val_accu
racy: 0.7981
Epoch 30/30
391/390 [=====] - 62s 158ms/step - loss: 0.5985 - accuracy: 0.7910 - val_loss: 0.6238 - val_accu
racy: 0.7957
```

Plot the model

```
In [12]: # Plot the training and validation accuracy and loss
accuracy = history.history["accuracy"]
val_accuracy = history.history["val_accuracy"]
loss = history.history["loss"]
val_loss = history.history["val_loss"]
epochs = range(1, len(accuracy) + 1)
plt.plot(epochs, accuracy, "bo", label="Training accuracy")
plt.plot(epochs, val_accuracy, "b", label="Validation accuracy")
plt.title("Training and validation accuracy")
plt.legend()
plt.figure()
plt.plot(epochs, loss, "bo", label="Training loss")
plt.plot(epochs, val_loss, "b", label="Validation loss")
plt.title("Training and validation loss")
plt.legend()
plt.show()
```



Model Evaluation

In [13]:

```
# Evaluate the convnet
test_loss, test_acc = model.evaluate(x_test, y_test)
print(f'Test accuracy: {test_acc*100:.1f}%')
print(f'Test loss: {test_loss:.3f}')
```

```
313/313 [=====] - 4s 13ms/step - loss: 0.6238 - accuracy: 0.7957
Test accuracy: 79.6%
Test loss: 0.624
```

The accuracy score with dropout and data augmentation has been increased to 79.6%

Save the model

```
In [14]: model.save('results/with_dropout_augmentation')

WARNING:tensorflow:From /opt/conda/lib/python3.8/site-packages/tensorflow/python/ops/resource_variable_ops.py:1813: calling BaseResourceVariable.__init__ (from tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and will be removed in a future version.
Instructions for updating:
If using Keras pass *_constraint arguments to layers.
INFO:tensorflow:Assets written to: results/with_dropout_augmentation/assets
```

Assignment 6.3

Load the ResNet50 model. Perform image classification on five to ten images of your choice. They can be personal images or publically available images. Include the images in dsc650/assignments/assignment06/images/. Save the predictions dsc650/assignments/assignment06/results/predictions/resnet50 directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

Loading libraries

```
In [7]: # Load libraries
from tensorflow.keras.applications.resnet50 import ResNet50
from tensorflow.keras.preprocessing import image
from tensorflow.keras.applications.resnet50 import preprocess_input, decode_predictions
import numpy as np
from IPython.display import Image, display
import os
```

Define Model

```
In [2]: # Load model
model = ResNet50(weights='imagenet')
```

Image Classification

```
In [13]: ## Custom function to predict the input image using resnet50
def image_prediction(img_input):
    ## model prediction and printing result
```

```

img_path = img_input
img = image.load_img(img_path, target_size=(224, 224))
x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)

preds = model.predict(x)

# decode the results into a list of tuples (class, description, probability)
print("Displaying the prediction result for the image: {}".format(image))
print('Predicted:', decode_predictions(preds, top=3)[0])

```


In [16]:

```

## Reading the image present in images directory
## calling image_prediction function
for img in os.listdir('images'):
    input_img = "images/"+img
    print('\nDisplaying {} image'.format(img))
    dis = Image(filename=input_img)
    display(dis)
    image_prediction(input_img)

```

Displaying deer.jpg image



Displaying the prediction result for the image: <module 'tensorflow.keras.preprocessing.image' from '/opt/conda/lib/python3.8/site-packages/tensorflow.keras/preprocessing/image/__init__.py'>
 Predicted: [('n12998815', 'agaric', 0.09936827), ('n02423022', 'gazelle', 0.08818725), ('n02115913', 'dhole', 0.079387225)]

Displaying polar.jpg image



Displaying the prediction result for the image: <module 'tensorflow.keras.preprocessing.image' from '/opt/conda/lib/python3.8/site-packages/tensorflow.keras/preprocessing/image/__init__.py'>

Predicted: [('n02510455', 'giant_panda', 0.99944896), ('n02447366', 'badger', 0.00021097742), ('n02134084', 'ice_bear', 0.0001567416)]

Displaying hipo.jpg image



Displaying the prediction result for the image: <module 'tensorflow.keras.preprocessing.image' from '/opt/conda/lib/python3.8/site-packages/tensorflow.keras/preprocessing/image/__init__.py'>

Predicted: [('n02422106', 'hartebeest', 0.23387302), ('n02410509', 'bison', 0.15939221), ('n02132136', 'brown_bear', 0.06565593)]

Displaying dolphin.jpg image



Displaying the prediction result for the image: <module 'tensorflow.keras.preprocessing.image' from '/opt/conda/lib/python3.8/site-packages/tensorflow.keras/preprocessing/image/__init__.py'>

```
Predicted: [('n02071294', 'killer_whale', 0.878125), ('n01484850', 'great_white_shark', 0.07984153), ('n01491361', 'tiger_shark', 0.013457938)]
```

Displaying dog.jpg image



Displaying the prediction result for the image: <module 'tensorflow.keras.preprocessing.image' from '/opt/conda/lib/python3.8/site-packages/tensorflow.keras/preprocessing/image/__init__.py'>

```
Predicted: [('n02091635', 'otterhound', 0.464588), ('n02099601', 'golden_retriever', 0.24075353), ('n02113799', 'standard_poodle', 0.09137372)]
```

Displaying zebra.jpg image



Displaying the prediction result for the image: <module 'tensorflow.keras.preprocessing.image' from '/opt/conda/lib/python3.8/site-packages/tensorflow.keras/preprocessing/image/__init__.py'>

```
Predicted: [('n02391049', 'zebra', 0.9975666), ('n02422106', 'hartebeest', 0.001230741), ('n02422699', 'impala', 0.000562362)]
```

Displaying tiger.jpg image



Displaying the prediction result for the image: <module 'tensorflow.keras.preprocessing.image' from '/opt/conda/lib/python3.8/site-packages/tensorflow/keras/preprocessing/image/__init__.py'>
Predicted: [('n02129604', 'tiger', 0.87079287), ('n02123159', 'tiger_cat', 0.110966384), ('n02391049', 'zebra', 0.006612492)]

In []: